


Lecture 1 - Review of classical inference

Guiding example: testing claims of **therapeutic touch**

- ↳ google "Emily Rosa"

- ↳ practitioners claim to feel "Human Energy Field"

- ↳ Rosa's experiment: experimenter randomly selects one of the practitioner's hands and holds her hand above it. (all occluded behind screen)
Practitioner is asked to identify which hand the experimenter has selected.

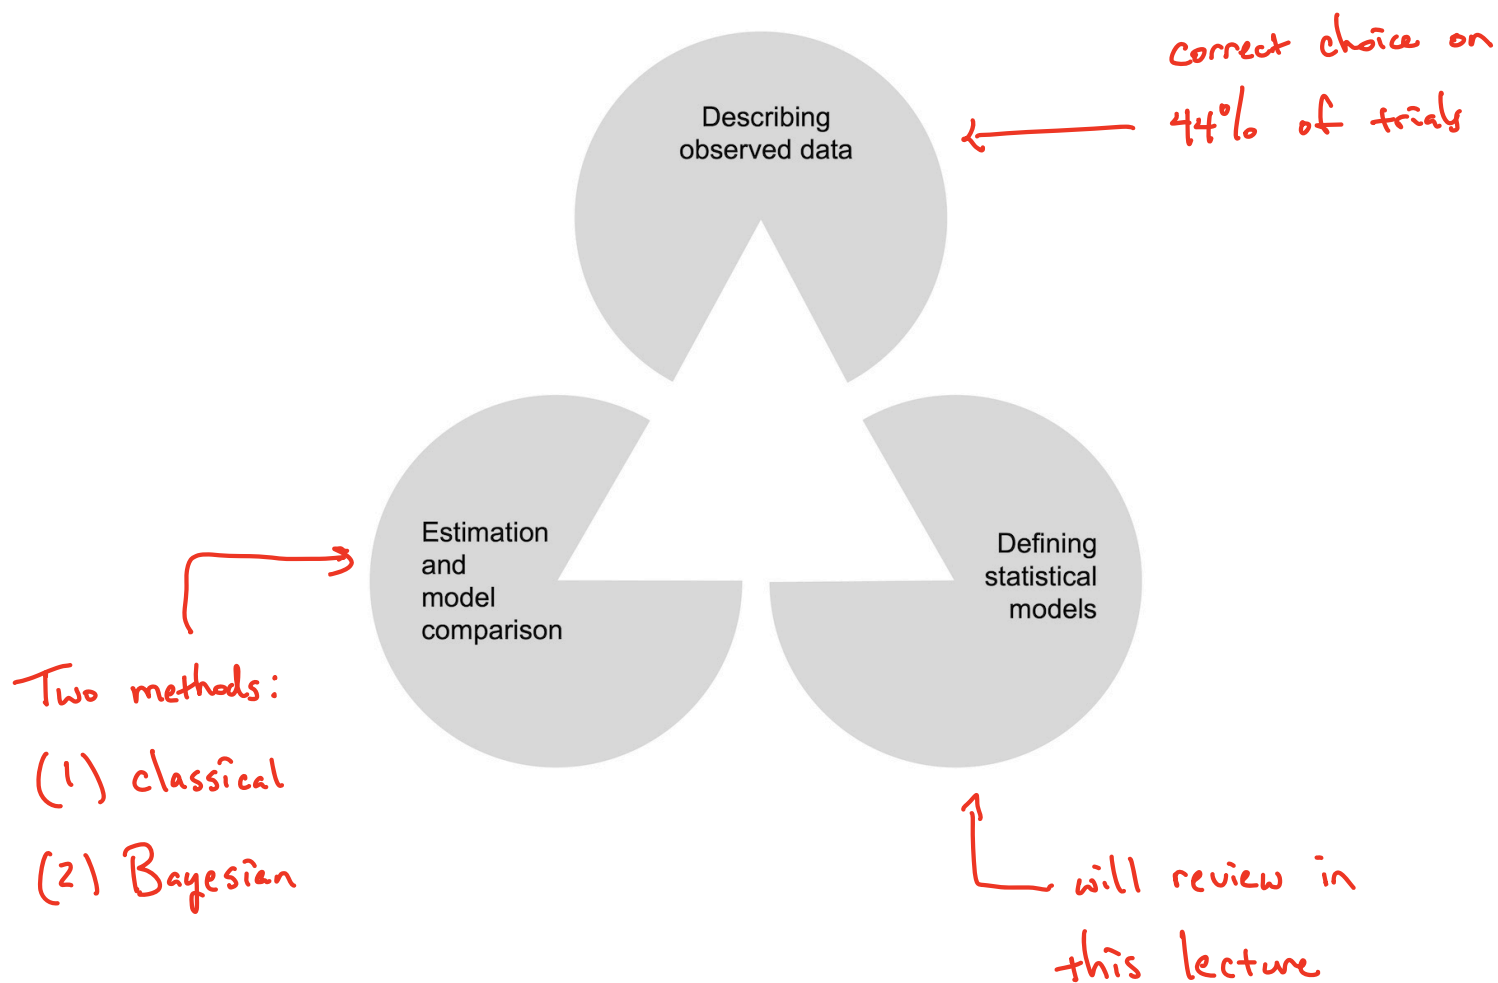


- ↳ Results: practitioners only correct on 44% of trials.

- ↳ Goal: **evaluate practitioners' claims**

Framework of inferential statistics

(from Faulkenberry, 2022)





What is a model?

↳ **mathematical function** that links some latent (population) variables to observable data.

→ probability distribution

Visualization:



Probability of
getting a success
on a trial (w)

Binomial model

Number of
successes out of
a set of trials (x)

Binomial model - for a given probability of success w , the binomial model gives probability of observing x successes across N trials.

Example: Suppose $N = 10$ trials, probability of success on each trial is $w = 0.25$. What is probability of observing $x = 8$ successes?

* 10 trials must be either Success or Failure.

↳ 8 Successes \rightarrow 2 Failures

Observe: S S S S S S S S F F

Prob: $\left[0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \cdot 0.25 \right] \cdot \left[0.75 \cdot 0.75 \right]$

$$= (0.25)^8 \cdot (0.75)^2$$

But, there are multiple ways to get 8 Successes & 2 Failures.

Ans: $\binom{10}{8} = \text{"10 choose 8"} = \frac{10!}{2! \cdot 8!}$

$$= \frac{10 \cdot 9 \cdot 8 \cdot 7 \cdot 6 \cdot 5 \cdot 4 \cdot 3 \cdot 2 \cdot 1}{2 \cdot 1 \cdot 8 \cdot 7 \cdot 6 \cdot 5 \cdot 4 \cdot 3 \cdot 2 \cdot 1} = 45$$

$$\begin{aligned}
 \text{So } P(X=8) &= \binom{10}{8} \cdot (0.25)^8 (0.75)^2 \\
 &= 0.000386 \\
 &\quad (\text{NOT very likely!})
 \end{aligned}$$


Can do these computations in R!

```

6 # Example: w = 0.25, N = 10 trials, probability of 8 successes?
7 choose(10,8) * (0.25)^8 * (0.75)^2
8
9 # R has a built-in function for this!
10 dbinom(x=8, size=10, prob=0.25)

```

How does the binomial model work as a model?




Start with a value
for w (prob of success
on each trial)

Generate potential data
that could be observed.

Plotting the potential data

```
13 x = seq(from=0, to=10, by=1)
14 barplot(dbinom(x, size=10, prob=0.25),
15         names.arg = 0:10,
16         xlab="Data x",
17         ylab="P(x)")
```


* small numbers of successes
more likely.




```
13 x = seq(from=0, to=10, by=1)
14 barplot(dbinom(x, size=10, prob=0.75),
15         names.arg = 0:10,
16         xlab="Data x",
17         ylab="P(x)")
```

← only change the
parameter w .

* large numbers of successes
more likely.



Problem of inference - given data, what is parameter?



What is the value of the parameter w ?

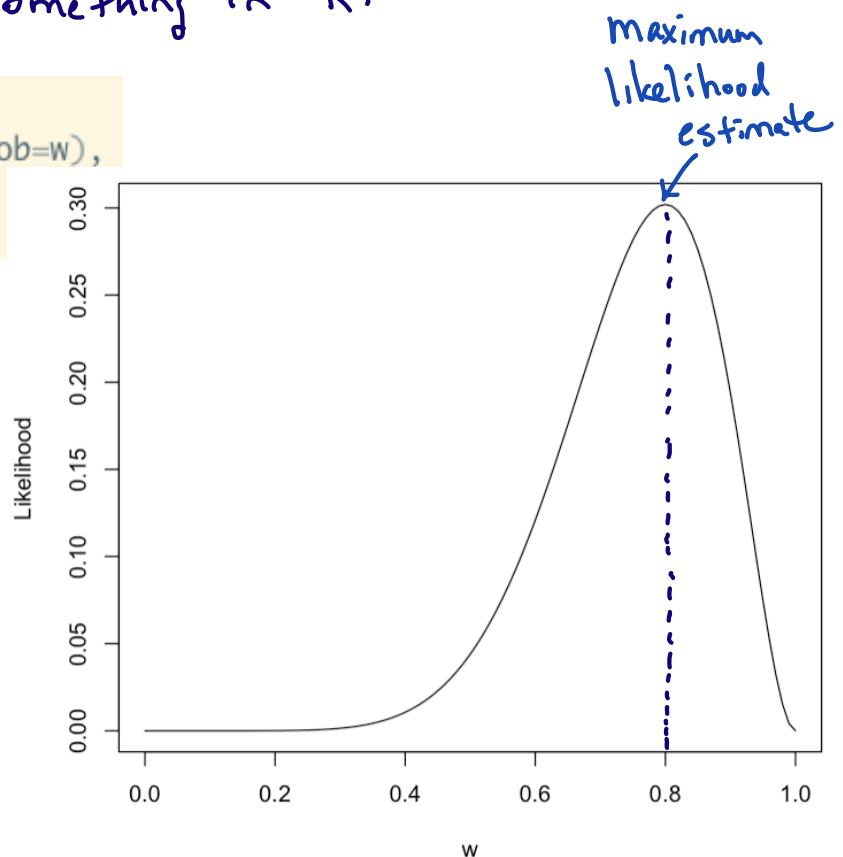


Suppose we observe $x=8$ successes in 10 trials.

To solve this, let's plot something in R:

```
25 w = seq(from=0, to=1, by=0.01)
26 plot(w, dbinom(x=8, size=10, prob=w),
27       type="l",
28       ylab="Likelihood")
```

"Likelihood function"



Problem of inference

↳ given data, what is/are the value(s) of the model parameter(s)?

↳ from above, we can see that there is uncertainty with respect to these parameter values

↳ Bayesian statistics provides a comprehensive way to quantify this uncertainty.

↳ more in Lecture 2.

As a review, let's talk about model comparison (hypothesis testing)

Step 1: set up two models / hypotheses

* Null hypothesis - performance based on "guessing"

$$\hookrightarrow H_0: w = 0.5$$

* Alternative hypothesis - performance better than guessing

$$\hookrightarrow H_1: w > 0.5$$

Step 2 - compare the models

↳ which one best predicts the observed data?

↳ classical method (due to R. Fisher, ~1920):

* Assume H_0 is true. Then calculate the probability of observing the data (or more extreme) under this assumption.


* compute $P(\text{data} \mid H_0)$

$$= P(x \geq 7 \mid \omega = 0.5)$$

= "p-value"

Let's use R:

```
45 x = seq(from=0, to=10, by=1)
46 barplot(dbinom(x, size=10, prob=0.5), # change to 0.5 to reflect null
47         names.arg = 0:10,
48         xlab="Data x",
49         ylab="P(x)")
```



Computing this p-value is easy in R

↳ "pbinom" function - computes $\text{prob} \leq \text{given data}$.

↳ since we want $\text{prob} \geq \text{data}$, must invert it and adjust
($1 - \text{pbinom}$)

```
53 1 - pbinom(7, size=10, prob=0.5)
```

= 0.0547

This p-value is small (i.e., $p = 0.05$) which means that the observed data is rare under H_0 .

Thus, we reject H_0 as a plausible model, leaving us with support for H_1 .

Note: that was a conceptual way to perform the hypothesis test. R will do it directly!

```
> binom.test(x=8, n=10, p=0.5, alternative="greater")
```

Exact binomial test

data: 8 and 10

number of successes = 8, number of trials = 10, p-value = 0.05469

alternative hypothesis: true probability of success is greater than 0.5

95 percent confidence interval:


0.4930987 1.0000000

sample estimates:

probability of success

0.8

Back to the Emily Rosa study:



↳ Results: practitioners only correct on 44% of trials.
(actual data from JAMA paper - 123 out of 280 trials).

$H_0: w = 0.5$ (practitioners were guessing)

$H_1: w > 0.5$ (practitioners performing above chance)

```
> binom.test(x=123, n=280, p=0.5, alternative="greater")
```

Exact binomial test

data: 123 and 280

number of successes = 123, number of trials = 280, p-value = 0.9819

alternative hypothesis: true probability of success is greater than 0.5

95 percent confidence interval:

0.3893703 1.0000000

sample estimates:

probability of success

0.4392857

observed data
very plausible
under H_0